Yuhao Ge

Mobile: (217)926-3291 | Email: yuhaoge2@illinois.edu | Homepage | LinkedIn

Education

University of Illinois at Urbana-Champaign M.Sc. Computer Science GPA: **4.0**/4.0 2023.8 - 2025.5

University of Illinois at Urbana-Champaign B.Sc. Computer Engineering GPA: **3.95**/4.0 2019.8 - 2023.5

- Honors: Highest Honors, Bronze Tablet (3%, 2023), Dean's List (2020&2022)

Zhejiang University B.Eng. Computer Engineering GPA: **3.93**/4.0 2019.8 - 2023.5

- Honors: Zhejiang Provincial Government Scholarship, First-Class Scholarship (3%, 2020&2022)

Skills

- **Programming:** Python, C/C++, JavaScript, SystemVerilog, Assembly, SQL, MongoDB, Neo4j
- **Frameworks & Tools:** CUDA, PyTorch, TensorFlow, TVM, Triton, PyG, Flask, React, Node.js, Kubernetes, AWS, GCP

Work Experience

Amazon, Annapurna Labs | *Software Engineer Intern* | *Compiler, Systems, ML, Accelerators* 2024.5 – 2024.8

- Neuron Compiler | *Optimize deep learning on AWS AI accelerators (Trainium and Inferentia)*
  - Delivered **three** organization-wide presentations and received a **Certificate of Appreciation** for developing one of the most influential internship projects, which significantly enhanced performance on Trainium and opened up a new direction for future development.
  - Developed infrastructure for automatic kernel generation, compilation, profiling, and visualization with a defined sweep space
  - Collected data for **DMA access pattern** analysis and introduced a **learning-based DMA latency model**
  - Created the **first-generation Autotuning** infrastructure from scratch for **compiler optimization** and **kernel optimization**
  - Used autotuning to optimize the Matrix Multiply **Fusion Pass**, achieving a 14.7% improvement for the **Llama3.1** model
  - Developed **kernel language** for AI accelerators and supported the **kernel optimization** with autotuning, achieving a 4.9% HFU improvement for kernels like Matrix Multiply
  - Implemented **multi-process compilation** and **distributed benchmarking**, resulting in 8.62X speedup

TikTok | *Software Engineer Intern* | *C++, Lua, Game Engine, AR/VR* 2022.5 - 2022.8

- Amazing Engine | *TikTok's Next-Generation 3D Game Engine for AR/VR Effects*
  - Collaborated in developing TikTok's **3D Game Engine**, which empowers users to create/use interactive **AR/VR** stickers
  - Implement a query-based animation system **Motion Matching** in **C++** for realistic and responsive avatar control
  - Developed an **SDK** for **Skeleton Retargeting** in **C++** and **Lua**, supporting animation adaptation across character models
  - Integrated the cross-functional team's **Text-to-Animation** algorithm into our game engine using the developed SDK

University of California, Los Angeles | *Visiting Student Researcher* | *ML, RL, GNN, FPGA, EDA* 2022.6 - 2022.11

- GNNDSE | *An automated design space exploration for automatic FPGA accelerator design* | *Advisor: Prof. Jason Cong*
  - Combined **GNN** with an **ML/RL**-based **Design Space Exploration** to achieve **FPGA** Accelerator **Design Automation**
  - Developed a learning-based **Cost Model** with **GNN** as a surrogate of the HLS tool for quick and accurate assessment
  - Optimized DSE by deploying heuristic algorithms such as **Genetic Algorithm** and **Simulated Annealing**
  - Used **Reinforcement-Learning** and **Bandits** for automatic algorithm selection, boosting exploration speed by 11%

Selected Projects

Optimized GPU code generation framework for SParse reguLar Attention 2023.8 – 2024.7

- Developed SPLAT, an optimized framework for efficient **sparse-MHSA**, targeting moderate sparsity levels
- Introduced **Affine Compressed Sparse-Row** (ACSR) format for **regular sparsity** patterns in MHSA
- Engineered advanced **GPU code-generation algorithms** for ACSR, enhancing sparse-MHSA kernel performance
- Achieved 2.05x and 4.05x speedups over **Triton** and **TVM** kernels with SPLAT implementation

GoPricing: An NFT pricing service powered by machine learning 2023.2 – 2023.6

- Developed a **Regression Model** for NFT pricing using historical transaction data and NFT features
- Design a system to support **online inference** and **offline training**, testing, and maintaining
- Implemented with **FastAPI**, **Airflow**, **MongoDB,** and **Redis**, deployed with **Kubernetes,** and monitored through **Grafana**

Remote Car Control System with Real-time 3D Reconstruction 2023.1 – 2023.5

- Developed a Raspberry Pi robot car with remote control via joysticks, utilizing PID control and STM32 microcontroller
- Implemented WiFi-based communication for transmission of commands and RGBD images between the car and server
- Implemented the **SLAM** framework RTAB-Map on the server for real-time **3D reconstruction**, achieving a 10Hz framerate

Implement A Game Efficiently on the FPGA Board 2022.1 - 2022.5

- Ported the game "Doodle Dump" to FPGA with **SystemVerilog**, achieving low power consumption and high efficiency
- Implemented a SOC with NIOS II in **C** to manage complex tasks like USB protocol and memory I/O
- Consumed only 400KB memory, 0.5w power to achieve a 50hz frame rate, won the **Best Design** Prize

Selected Research

- SPLAT: A framework for optimised GPU code-generation for SParse reguLar ATtention