

Yuhao Ge geyuhao33@gmail.com | (217) 926-3291 | [Homepage](#) | [Linkedin](#)

LLM serving and performance engineer with **cross-accelerator expertise** (TPU, Trainium, GPU), spanning **framework/compiler/kernel** optimizations in large-scale serving systems.

Education

University of Illinois at Urbana-Champaign | M.S. in Computer Science, 2023/08 - 2024/12

University of Illinois at Urbana-Champaign | B.S. in Computer Engineering, GPA 4.0/4.0 2019/08 - 2023/05

– Honors: Highest Honors, Bronze Tablet (3%, 2023), Dean's List (2020 & 2022)

Skills

ML Frameworks & Serving: PyTorch, JAX, vLLM, SGLang

Kernels & Compilers: Pallas, Triton, NKI, CUDA, XLA, MLIR, TVM

Experience

Google | Software Engineer | TPU, vLLM, PyTorch, Pallas, Gemini 2025/11 - Present

- **Led a team** to build the **PyTorch-native vLLM TPU backend from 0 to 1**, making TPU a drop-in target for vLLM users with GPU-equivalent usability and performance, now serving key customers' production workloads.
- Designed and implemented the whole **TPU serving stack from scratch**: async scheduling, eager and compiled mode, static-shape bucketing and padding strategy, distributed serving across TP/EP/PP/DP, and Pallas kernel integration.
- Co-developed **torchtpu**, the PyTorch-TPU bridge underlying the backend.
- Brought up and optimized major OSS models (**GPT-OSS, Qwen3, Qwen3.5**) across **framework, kernel, and compiler** layers, reaching 80% of GB300 performance on TPuv7.
- Design TPU-friendly **hybrid cache management system** for hybrid models like Qwen3.5 and Kimi K3.
- Designed and developed TPU-friendly **quantization** schemes (**W4A8, W8A8, KV cache**) and checkpoint support for **NVFP4, INT4, FP8, and MXFP4** formats.
- Drove OSS co-development with **Inferact** and **RadixArk**, designing tiered **CI/CD** (presubmit, postsubmit, nightly) on TPU hardware with **accuracy and performance regression gates** for cross-org contribution and release health.

AWS, Annapurna Labs | Software Engineer | ML Compiler, LLMs, LLM Systems, Architecture 2025/03 - 2025/11

- **Founding engineer** of the **NKI** kernel language and compiler for AWS **Trainium**, driving language and compiler design.
- Architected and optimized core kernel compiler passes (**scheduling, allocation**) with **cycle-level optimization**.
- Designed and implemented custom **FlashAttention** kernels tailored for LLM workloads on Trainium.
- Led development of **autotuning** infrastructure to systematically optimize compiler and kernel performance.

AWS, Annapurna Labs | Software Engineer Intern | ML Compiler, LLM Systems 2024/05 - 2024/08

- Developed automated **kernel generation** and profiling infrastructure to collect micro-architectural performance data and train **ML-based cost models** for instruction latency estimation.
- Architected the **first-generation autotuning** infrastructure from scratch, designing a modular, pass-agnostic framework with clean separation of design space, search strategy, and surrogate cost modeling.
- Used autotuning to optimize the Matmul Fusion Pass, achieving a 14.7% improvement for the Llama3.1 model.
- Implemented multi-process compilation and distributed profiling, accelerating autotuning workflows by 8.6×.
- Recognized with a [Certificate of Appreciation](#) for high-impact contributions.

TikTok | Software Engineer Intern | ML, Graphics, Game Engine, AR/VR. 2022/05 - 2022/08

- Developed TikTok's 3D Game Engine, which empowers users to create/use interactive AR/VR stickers.
- Implemented a query-based Motion Matching system in C++ for real-time, low-latency avatar control.
- Developed an SDK for Skeleton Retargeting, supporting animation adaptation across character models

University of California, Los Angeles | Research Assistant | ML/RL, FPGA 2022/06 - 2022/11

- Developed a **GNN-based surrogate cost model** to approximate HLS scheduling and resource estimation.
- Designed an ML-driven design space exploration framework for automated FPGA accelerator optimization.
- Integrated **reinforcement learning** and bandit strategies for adaptive search, improving exploration efficiency by 11%.

Publications

- [SPLAT](#): A framework for optimised GPU code-generation for SParse reguLar ATtention (OOPSLA' 25).

Projects

Optimized GPU code generation framework for SParse reguLar Attention 2023/08 - 2024/07

- Developed SPLAT, an optimized framework for efficient sparse-MHSA, targeting moderate sparsity levels.
- Introduced Affine Compressed Sparse-Row (ACSR) format for regular sparsity patterns in MHSA.
- Engineered advanced GPU code-generation algorithms for ACSR, enhancing sparse-MHSA kernel performance.
- Achieved 2.05x and 4.05x speedups over Triton and TVM. Accepted at OOPSLA 2025.

Implement A [“Doodle Jump”](#) Efficiently on the FPGA Board

2022/01 - 2022/05

- Ported the game “Doodle Jump” to FPGA with SystemVerilog, achieving low power consumption and high efficiency.
- Implemented a SOC with NIOS II in C to manage complex tasks like USB protocol and memory I/O.
- Consumed only 400KB memory, 0.5w power to achieve a 50 HZ frame rate, won the *Best Design Prize*.