

Yuhao Ge

Mobile: (217)926-3291 | Email: yuhao2@illinois.edu | [Homepage](#) | [LinkedIn](#)

Education

University of Illinois at Urbana-Champaign	Master of Computer Science		2023.8 – 2025.1
University of Illinois at Urbana-Champaign	B.Sc. Computer Engineering	GPA: 4.0/4.0	2019.8 - 2023.5
■ Honors: Highest Honors, Bronze Tablet (3%, 2023), Dean's List (2020&2022)			
Zhejiang University	B.Eng. Computer Engineering	GPA: 3.93/4.0	2019.8 - 2023.5
■ Honors: Zhejiang Provincial Government Scholarship, First-Class Scholarship (3%, 2020&2022)			

Skills

- **Programming:** Python, C/C++, JavaScript, SystemVerilog, Assembly, SQL, MongoDB, Neo4j
- **Frameworks & Tools:** CUDA, PyTorch, TensorFlow, XLA, TVM, Triton, PyG, Flask, React, Node.js, Kubernetes, AWS, GCP

Work Experience

Amazon, Annapurna Labs <i>Software Engineer Compiler, Systems, ML, Accelerators</i>	2025.3 – now
■ Neuron Compiler <i>Optimize deep learning on AWS AI accelerators (Trainium and Inferentia)</i> • Developed a loop software pipelining pass supporting user-guided scheduling and deterministic transformations. • Implemented modulo scheduling with a cost model for compiler-driven optimization and added autotuning for continuous performance improvement. • Designed and extended the Neuron Kernel Interface (NKI) to provide fine-grained user control and improved usability. • Integrated user-configurable backend passes (e.g., vectorization, loop scheduling) into the kernel compilation flow	
Amazon, Annapurna Labs <i>Software Engineer Intern Compiler, Systems, ML, Accelerators</i>	2024.5 – 2024.8
■ Neuron Compiler <i>Optimize deep learning on AWS AI accelerators (Trainium and Inferentia)</i> • Delivered three organization-wide presentations and received a Certificate of Appreciation for developing the most influential internship projects, significantly enhancing performance, and opening up new directions. • Developed infrastructure for automatic kernel generation, compilation, profiling, and visualization with a defined sweep space • Collected data for DMA access pattern analysis and introduced a learning-based DMA latency model • Created the first-generation Autotuning infrastructure from scratch for compiler optimization and kernel optimization • Used autotuning to optimize the Matrix Multiply Fusion Pass, achieving a 14.7% improvement for the Llama3.1 model • Designed and extended the Neuron Kernel Interface (NKI) to support the kernel optimization with autotuning, achieving a 4.9% HFU improvement for kernels like Matrix Multiply • Implemented multi-process compilation and distributed benchmarking, resulting in 8.62X speedup	
TikTok <i>Software Engineer Intern C++, Lua, Game Engine, AR/VR</i>	2022.5 - 2022.8
■ Amazing Engine <i>TikTok's Next-Generation 3D Game Engine for AR/VR Effects</i> • Collaborated in developing TikTok's 3D Game Engine, which empowers users to create/use interactive AR/VR stickers • Implement a query-based animation system Motion Matching in C++ for realistic and responsive avatar control • Developed an SDK for Skeleton Retargeting in C++ and Lua, supporting animation adaptation across character models • Integrated the cross-functional team's Text-to-Animation algorithm into our game engine using the developed SDK	
University of California, Los Angeles <i>Visiting Student Researcher ML, RL, GNN, FPGA, EDA</i>	2022.6 - 2022.11
■ GNNDSE <i>An automated design space exploration for automatic FPGA accelerator design Advisor: Prof. Jason Cong</i> • Combined GNN with an ML/RL-based Design Space Exploration to achieve FPGA Accelerator Design Automation • Developed a learning-based Cost Model with GNN as a surrogate of the HLS tool for quick and accurate assessment • Optimized DSE by deploying heuristic algorithms such as Genetic Algorithm and Simulated Annealing • Used Reinforcement-Learning and Bandits for automatic algorithm selection, boosting exploration speed by 11%	

Selected Projects

Optimized GPU code generation framework for SParse reguLar Attention	2023.8 – 2024.7
■ Developed SPLAT, an optimized framework for efficient sparse-MHSA , targeting moderate sparsity levels • Introduced Affine Compressed Sparse-Row (ACSR) format for regular sparsity patterns in MHSA • Engineered advanced GPU code-generation algorithms for ACSR, enhancing sparse-MHSA kernel performance • Achieved 2.05x and 4.05x speedups over Triton and TVM kernels with SPLAT implementation	
Implement A Game Efficiently on the FPGA Board	2022.1 - 2022.5
■ Ported the game "Doodle Dump" to FPGA with SystemVerilog , achieving low power consumption and high efficiency • Implemented a SOC with NIOS II in C to manage complex tasks like USB protocol and memory I/O • Consumed only 400KB memory, 0.5w power to achieve a 50hz frame rate, won the Best Design Prize	

Selected Research

- [SPLAT](#): A framework for optimised GPU code-generation for SParse reguLar ATtention (OOPSLA' 25)